



# Public Opinion Analysis of Airlines

---

## Team 1:

Fan Wu

Dayu Jia

Bowen Jiang

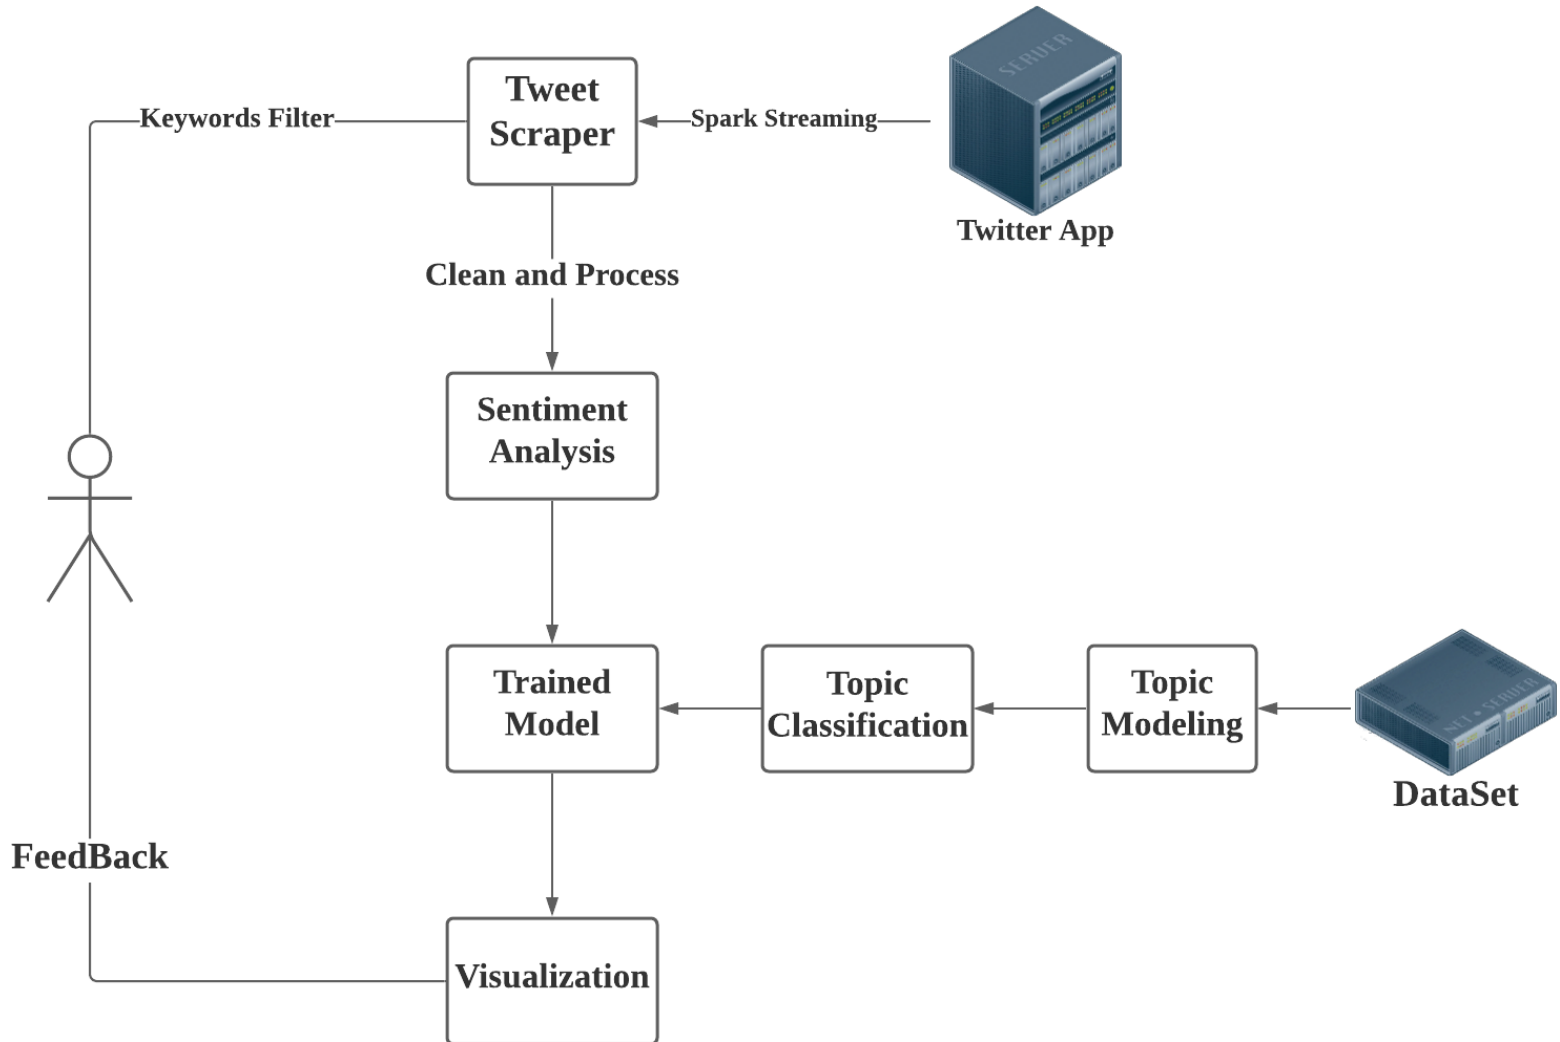
## Github link:

<https://github.com/FanWu6/Spark-stream-twitter-analysis>



# Use Cases

- User inputs the name of an airplane company
- and receives a comprehensive analysis on public opinion



# Methodology

- Tweets pulling -> Spark Streaming, twitter4j
- Data process -> StopWordsRemover, tf-idf, SparkSQL, etc...
- Topic modeling -> Latent Dirichlet Allocation(LDA)
- Topic classification -> Logistic Regression
- Sentiment analysis separately -> CoreNLP
- Visualization -> Elasticsearch, Kibana

# Data Sources

kaggle



## Twitter US Airline Sentiment

This dataset has 14485 rows and 30 columns

Source from:

<https://www.kaggle.com/crowdflower/twitter-airline-sentiment>

## Real-time data through Twitter API

Approximately 150 tweets in 1 minute.

(Depend on keywords put in)

# Milestones

An American Airlines airplane is shown in flight against a cloudy sky. The tail of the plane features the airline's signature red, white, and blue striped design. The word "American" is visible on the side of the fuselage.

## 1st week:

Implement two ways to extract real-time data from Twitter; Choose LDA to do tweets topic modeling.

## 2nd week:

Implement data cleaning and LDA; Learn about topic classification algorithm and sentiment analysis.

## 3rd week:

Implement Logistic Regression and train the model; Sentiment analysis real-time tweets; Learn about Elasticsearch.


## 4th week:

Uploading data to Kibana and visualization; Make our code better.

# Programming in Scala



Data extraction



Data processing



Model implement



Unit tests



Data uploading



# Acceptance criteria

- 70% of data can be classified correctly
- The accuracy of sentimental analysis should reach 80%

```
val evaluator = new BinaryClassificationEvaluator().setMetricName("areaUnderROC")
println("The topic prediction accuracy of test data = " + (evaluator.evaluate(predictions)))
//
LogisticRegressionModel
LogisticRegressionModel x
21/04/22 21:18:15 INFO SparkContext: Invoking stop() from shutdown hook
21/04/22 21:18:15 INFO SparkUI: Stopped Spark web UI at http://wufan-PC.lan:4040
The topic prediction accuracy of test data = 0.6761062958812633
```

```
val accuracy: Double = (accp * np + accn * nn + acco * no) / (nn + no + np)
println("Accuracy of Sentiment: " + accuracy)
SentimentAccuracy
SentimentAccuracy x
21/04/22 21:22:38 INFO SparkContext: Created broadcast 3 from rdd at Sentimen
21/04/22 21:22:38 INFO FileSourceScanExec: Planning scan with bin packing, ma
Accuracy of Sentiment: 0.9167457333333332
```

# Unit tests

```
Preprocess success.
[info] PreprocessSpec:
[info] Preprocess
[info] - should Preprocess work
[info] Run completed in 37 seconds, 62 milliseconds.
[info] Total number of tests run: 15
[info] Suites: completed 6, aborted 0
[info] Tests: succeeded 15, failed 0, canceled 0, ignored 0, pending 0
[info] All tests passed.
[success] Total time: 57 s, completed 2021年4月21日 7:39:54
21/04/21 19:39:54 INFO ContextCleaner: Cleaned accumulator 5213
21/04/21 19:39:54 INFO ContextCleaner: Cleaned accumulator 5205
```

# Goals

---

- For any airline company, we can offer a real-time analysis on public opinion based on our model.
- For us, we want to learn the using of Scala and machine learning, and how to co-work on Github.



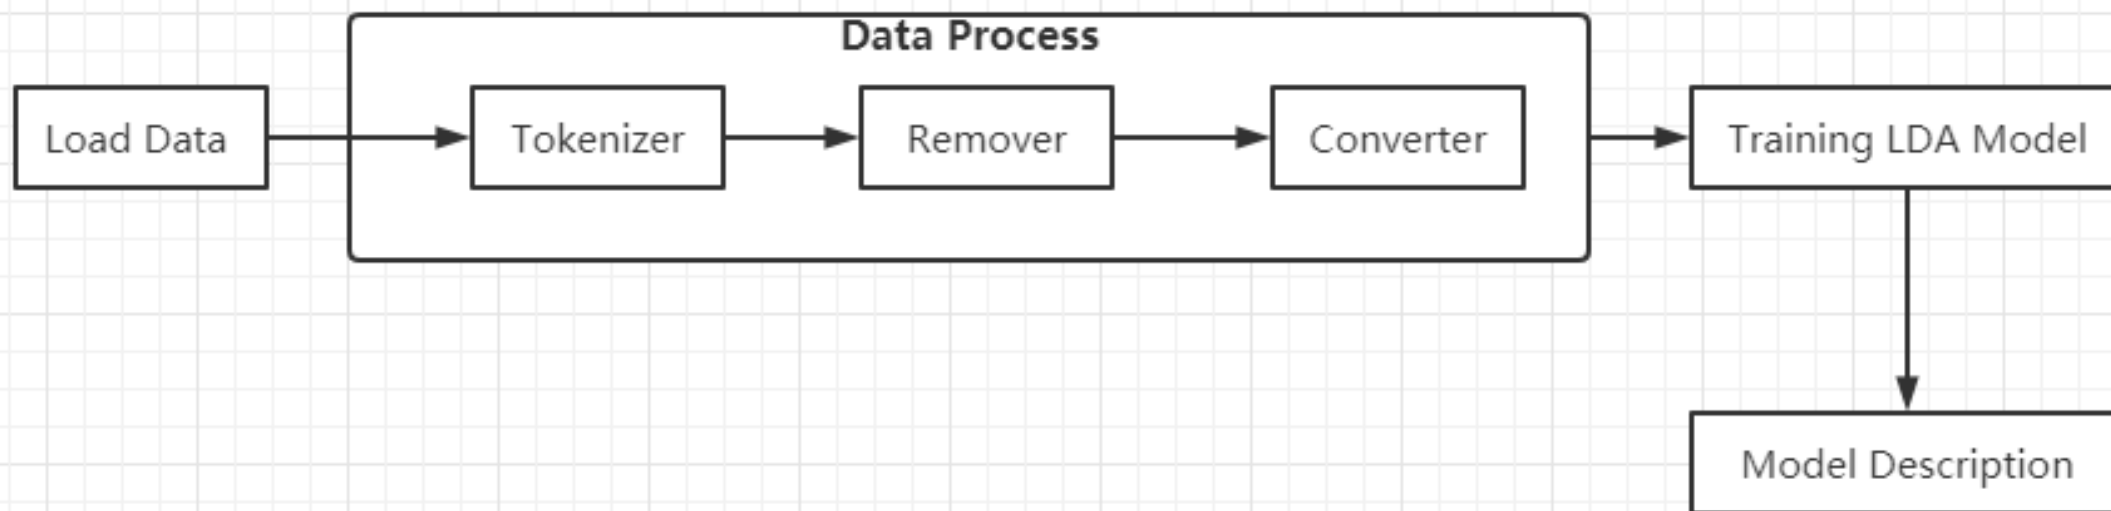


# Latent Dirichlet Allocation

---



# LDA Process Description



# Data Preprocess Code Description

```
//Tokenizing using the RegexTokenizer
val tokenizer = new RegexTokenizer()
  .setPattern("[\\W_]+")
  .setMinTokenLength(4)
  .setInputCol("corpus")
  .setOutputCol("tokens")

val tokenized_df: DataFrame = tokenizer.transform(corpus_df)

tokenized_df.select("tokens").show(10,false)

//Removing the Stop-words using the Stop Words remover
val add_stopwords = Array("http","jetblue","southwestair","americanair","flight",
  "usairways","thanks","virginamerica","thank","today","flightled","united","please")
val stopwords = sparkSession.read.text("data/actualdata/stopwords.txt")
  .collect().map(row => row.getString(0)).union(add_stopwords)

val remover = new StopWordsRemover()
  .setStopWords(stopwords) // This parameter is optional
  .setInputCol("tokens")
  .setOutputCol("filtered")

val filtered_df: DataFrame = remover.transform(tokenized_df)

//Converting the Tokens into the CountVector
val vectorizer: CountVectorizerModel = new CountVectorizer()
  .setInputCol("filtered")
  .setOutputCol("features")
  .setVocabSize(10000)
  .setMinDF(5)
  .fit(filtered_df)

val countVectors: DataFrame= vectorizer.transform(filtered_df).select("id", "features")
countVectors.show(5,false)
import sparkSession.implicitly._
val lda_countVector: RDD[(Long, linalg.Vector)] = countVectors.rdd.map {
  case Row(id: Long, countVector: Vector) => (id, Vectors.fromML(countVector)) }
(lda_countVector,vectorizer)
```

Using tokenizer filter away tokens with length <4

```
+-----+
|tokens|
+-----+
[[virginamerica, what, dhepburn, said]
[[virginamerica, plus, added, commercials, experience, tacky]
[[virginamerica, didn, today, must, mean, need, take, another, trip]
[[virginamerica, really, aggressive, blast, obnoxious, entertainment, your, guests, faces, they, have, little, recourse]
[[virginamerica, really, thing, about]
[[virginamerica, seriously, would, flight, seats, that, didn, have, this, playing]
[]
[[virginamerica, nearly, every, time, this, worm, away]
[[virginamerica, really, missed, prime, opportunity, without, hats, parody, there, https, mmpg7grezp]
[[virginamerica, well, didn]
+-----+
only showing top 10 rows
```

Remove stopwords

Set vocabulary size to 10000, and let the word show 5 times in each vocabulary

```
+-----+
|id|features|
+-----+
[0] |(2431,[92],[1.0])|
[1] |(2431,[74,331,839,1998],[1.0,1.0,1.0,1.0])|
[2] |(2431,[10,45,48,69,76,342,528],[1.0,1.0,1.0,1.0,1.0,1.0,1.0])|
[3] |(2431,[37,288,698,2155],[1.0,1.0,1.0,1.0])|
[4] |(2431,[37,282],[1.0,1.0])|
+-----+
only showing top 5 rows
```

Completes the process of converting the documents into a vector of word counts

# LDA Code Description

```
val sparkSession = SparkSession.builder()
  .appName( name = "LDA topic modeling")
  .master( master = "local[*]").getOrCreate()

val df: DataFrame = sparkSession.read.format( source = "csv")
  .option("header", "true")
  .load( path = "data/actualdata/Tweets.csv")

val processeddata = Preprocess.run(df, sparkSession)
val lda_countVector = processeddata._1
```

Load Data.

```
val lda = new LDA()
  .setOptimizer(new OnlineLDAOptimizer().setMiniBatchFraction(0.8))
  .setOptimizer("em")
  .setK(3)
  .setMaxIterations(100)
  .setDocConcentration(-1) // use default values
  .setTopicConcentration(-1) // use default values

val ldaModel: LDAModel = lda.run(lda_countVector)

//-----2.Model And Description-----
//Describe the final pre-maxtermspertopic words (the most important wo
val topicIndices = ldaModel.describeTopics(maxTermsPerTopic = 10)
val vocabList = processeddata._2.vocabulary
val topics = topicIndices.map { case (terms, termWeights) =>
  terms.map(vocabList(_)).zip(termWeights)
}

println(s"$numTopics topics:")
topics.zipWithIndex.foreach { case (topic, i) =>
  println(s"TOPIC $i")
  topic.foreach { case (term, weight) => println(s"$term\t$weight") }
  println(s"=====")
}
```

Builds LDA Model generate  
3 topic after 100 iterations.

```
TOPIC 2
plane    0.016412672845225632
cancelled 0.014447211525821221
gate     0.014399853550372406
delayed  0.013200672507164314
hours    0.011028695911745066
hour      0.01061138210683631
time      0.009173256673966691
late      0.009037141195921831
flights  0.0087523984615247
waiting  0.007621818729309813
=====
```

```
TOPIC 0
service  0.016767524916125754
customer 0.013806351161999138
cancelled 0.012375795622015757
hold      0.01198256160036785
time      0.011001898985795212
help      0.01042591942019184
hours     0.009209733292761754
still     0.009056049201819651
call      0.008279447544858572
phone     0.007334837209351291
=====
```

```
TOPIC 1
help      0.01276614950428537
cancelled 0.010189066432353023
service  0.00972013817244688
need      0.008325936422038775
customer 0.00825431325796964
time      0.00742962418454725
flights  0.007136676573303931
problems  0.006718731811876184
number    0.006643656326519289
change    0.006190771115745823
=====
```

# Some Optimizations To The Model

---

```
val lda = new LDA()  
  .setOptimizer(new OnlineLDAOptimizer().setMiniBatchFraction(0.8))  
  .setOptimizer("em")  
  .setK(4)  
  .setMaxIterations(100)  
  .setDocConcentration(-1) // use default values  
  .setTopicConcentration(-1) // use default values
```

- More iterations make the model more accurate
- Add new Stopwords to help filter

```
val add_stopwords = Array("http", "jetblue", "southwestair", "americanair", "flight",  
  "usairways", "thanks", "virginamerica", "thank", "today", "flightled", "united", "please")  
val stopwords = sparkSession.read.text("data/actualdata/stopwords.txt")  
  .collect().map(row => row.getString(0)).union(add_stopwords)
```





# Topic Classification

---





# Topic classification

---

For instance,

“We have the **gold level plan** and use it for everything, **love the features**! It is one of the **best bang for buck** possible.”

A topic classification model that’s been trained to understand these expressions (**gold level plan**, **love the features**, and **best bang for buck**) would be able to tag this review as topic of *Features* and *Price*.

In our case,

@VirginAmerica Is it me, or is your website down? BTW, your new website isn't a great **user experience**. Time for another redesign.

**Customer Service Issue**

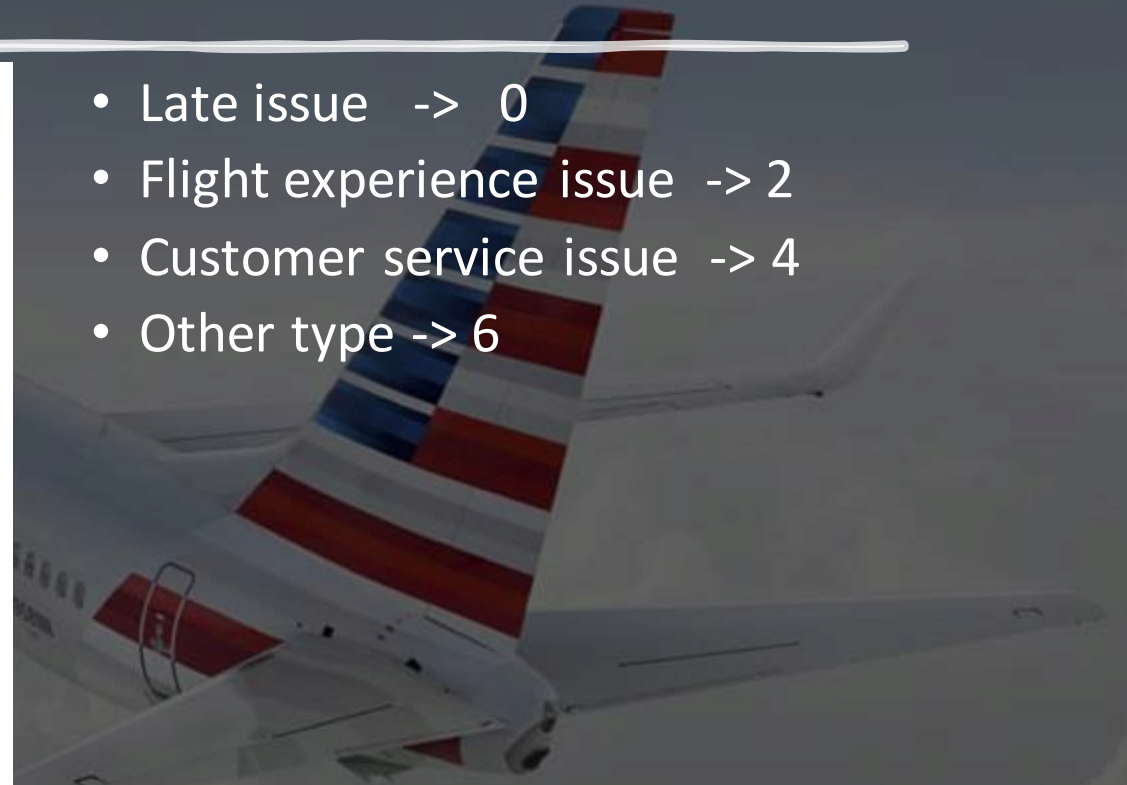
# Data for training

```
flights 0.0167903690142048
cancelled 0.013848024722761337
hours 0.010398252437989924
help 0.009549276062319553
still 0.008408700272805191
late 0.00823532559798192
delayed 0.007862819032489875
hour 0.007674428383153015
like 0.007281068289706297
email 0.006798501598634885
=====
TOPIC 1
gate 0.014362314033549274
service 0.011501779434624122
time 0.01147332879552762
plane 0.01051470294529229
call 0.009846823449676802
cancelled 0.00818193145594124
hours 0.008120184808420928
hold 0.00796607052723924
help 0.007811817503609371
customer 0.007377330252671869
=====
TOPIC 2
```

late issue

customer service  
issue

- Late issue -> 0
- Flight experience issue -> 2
- Customer service issue -> 4
- Other type -> 6



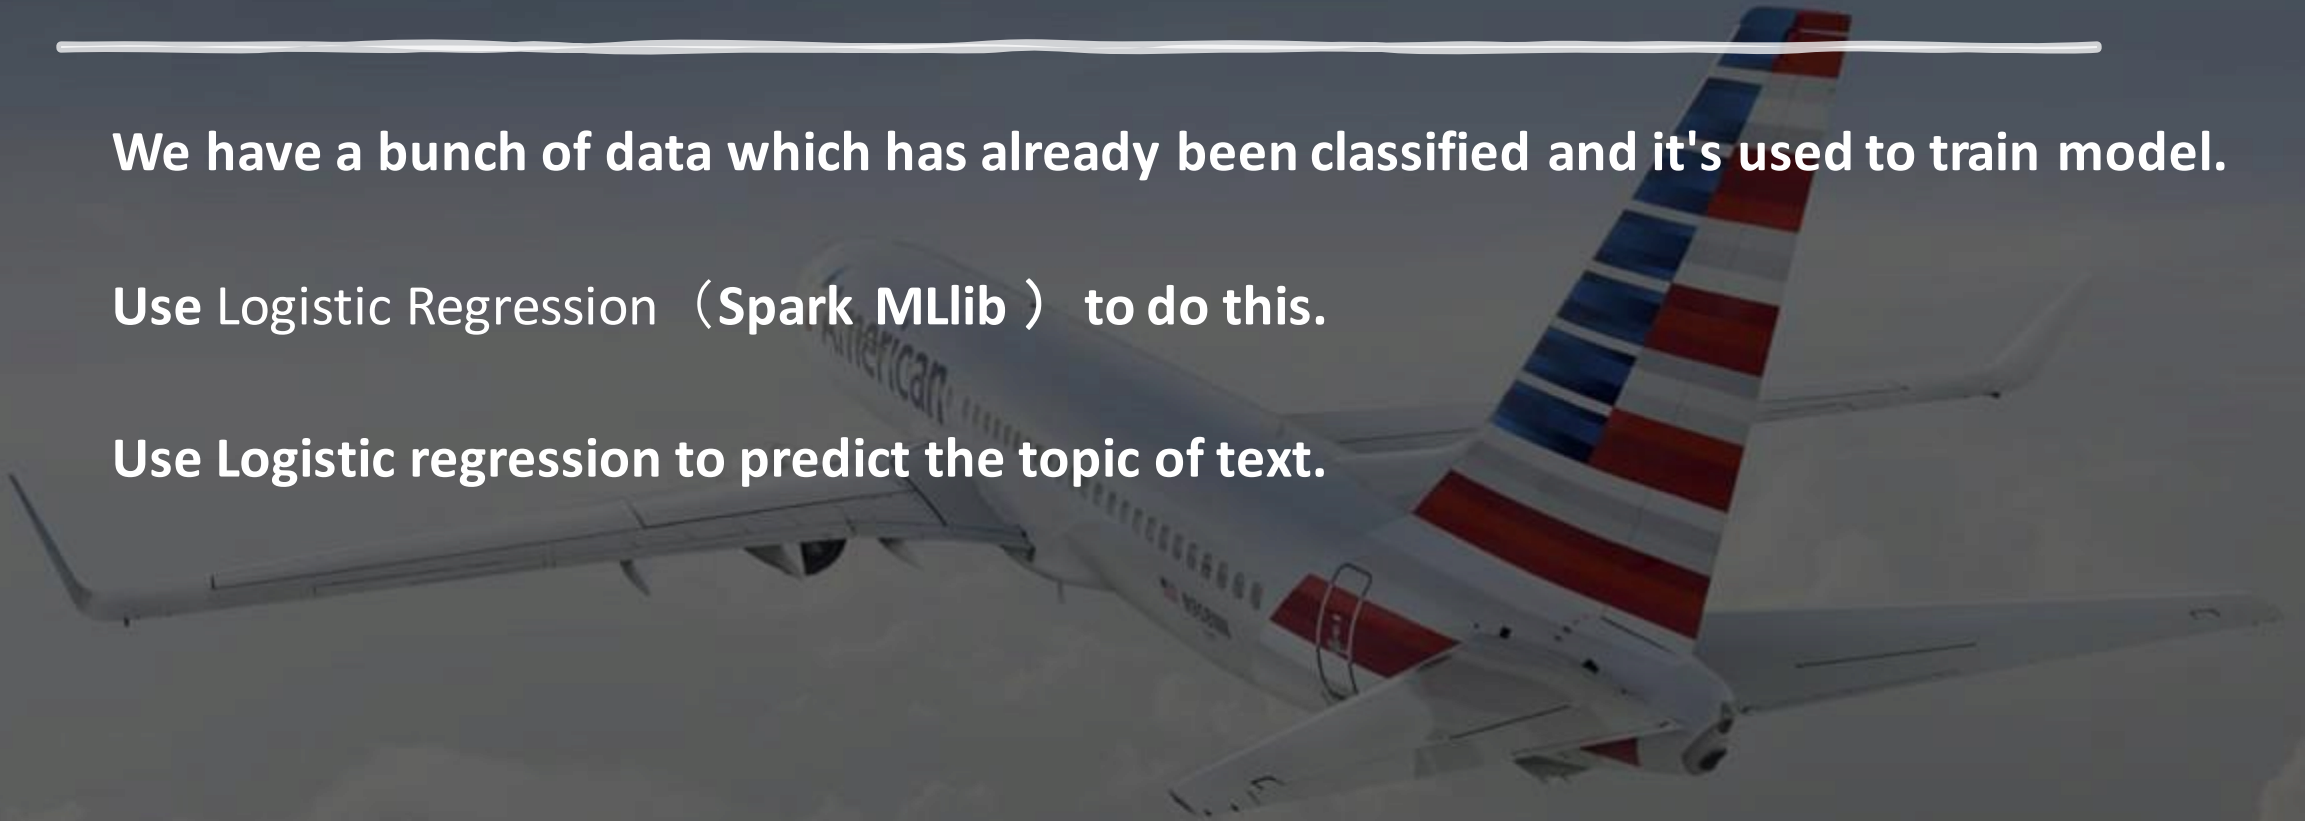
# Logistic Regression

---

**We have a bunch of data which has already been classified and it's used to train model.**

**Use Logistic Regression (Spark MLlib ) to do this.**

**Use Logistic regression to predict the topic of text.**



# Pre-process

```
//load training data
val newsgroupsRawData: RDD[String] = sc.textFile( path = "data/actualdata/train.csv")
//counts
println("The number of documents read in is " + newsgroupsRawData.count() + ".")

case class newsgroupsCaseClass(text: String, topic: String)

//remove all other character but words
val newsgroups: DataFrame = newsgroupsRawData.map{case (lines) =>
  val topic = lines.split( regex = "," ).take(1)(0)
  val text = TrainingUtils.processText(lines)
  newsgroupsCaseClass(text,topic)}.toDF()
newsgroups.cache()

newsgroups.printSchema()

newsgroups.sample( withReplacement = false, fraction = 0.001, seed = 10L ).show( numRows = 10, truncate = false)

newsgroups.groupBy( col1 = "topic" ).count().show()

//transform to another dataframe
val labelednewsgroups = newsgroups.withColumn( colName = "label", newsgroups("topic").cast( to = "double" ))

labelednewsgroups.sample( withReplacement = false, fraction = 0.003, seed = 10L ).show( numRows = 5, truncate = false)

//Split documents from a list of (id, text, label) tuples
val Array(training, test) = labelednewsgroups.randomSplit(Array(0.9, 0.1), seed = 12345)
```

1. load data

2. extract column text, and remove all other character but only leave words

3. cast topic to double

4. split data to training and test

```
+-----+-----+
|text                                         |topic|
+-----+-----+
| he was trying to take stuff from the under the seat in front of him and bugging out throughout the flight. feel safe. |2 |
| this means within one week i will have filed 2 compensation complaints to your website |4 |
| my wife is trying to get a group of clients to their destination and just got disconnected after2 hours holding. Help. |4 |
```

```
+-----+-----+-----+
|text                                         |topic|label|
+-----+-----+-----+
| he was trying to take stuff from the under the seat in front of him and bugging out throughout the flight. feel safe. |2 |2.0 |
| this means within one week i will have filed 2 compensation complaints to your website |4 |4.0 |
```

# Main stages

---

- Tokenizer
- StopWordsRemover
- TF-IDF
- Logistic Regression

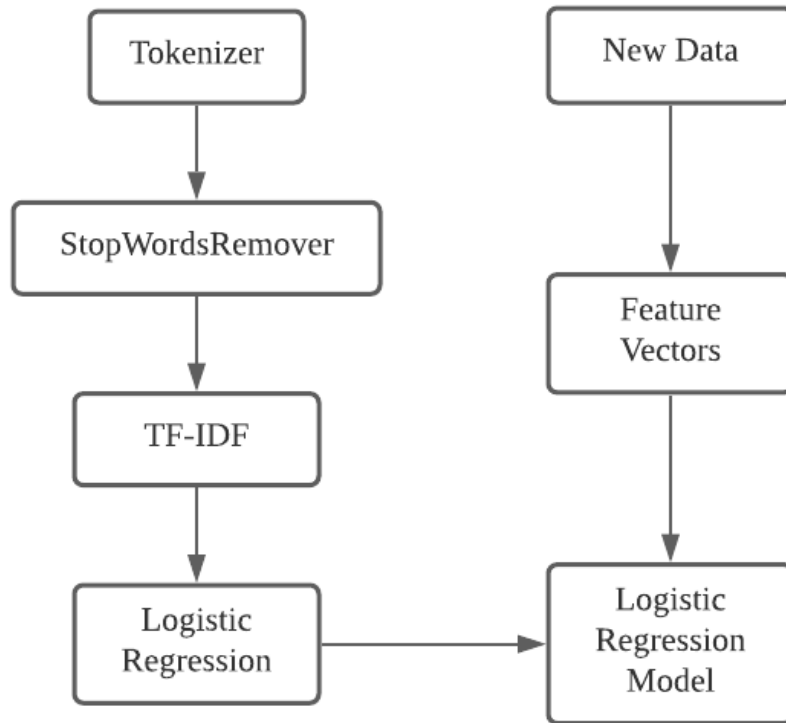
```
a  
about  
above  
after  
again  
against  
all  
am  
an  
and  
any  
are  
aren't  
arent  
as  
at  
be  
because
```



```
val tokenizer = new Tokenizer().setInputCol("text").setOutputCol("words")  
val remover = new StopWordsRemover().setInputCol("words").setOutputCol("filtered").setCaseSensitive(false)  
val hashingTF = new HashingTF().setNumFeatures(1000).setInputCol("filtered").setOutputCol("rawFeatures")  
val idf = new IDF().setInputCol("rawFeatures").setOutputCol("features").setMinDocFreq(0)  
val lr = new LogisticRegression().setRegParam(0.01).setThreshold(0.5)  
val pipeline = new Pipeline().setStages(Array(tokenizer, remover, hashingTF, idf, lr))
```

# Pipeline

---





# How to Optimize?

| text                | prediction | topic |
|---------------------|------------|-------|
| 2 and a half hou... | 0.0        | 4     |
| 2 hours on hold ... | 4.0        | 4     |
| I would love to ... | 6.0        | 2     |
| Im just praying ... | 0.0        | 2     |
| Instructions say... | 4.0        | 4     |
| Third flight in ... | 6.0        | 2     |
| Why offer automa... | 4.0        | 4     |
| Why you released... | 4.0        | 4     |
| Yo yo yo stuck o... | 0.0        | 0     |
| You respond to m... | 4.0        | 4     |
| airport and 2 ex... | 4.0        | 4     |
| anyone there to ... | 4.0        | 4     |
| but not sufficie... | 6.0        | 4     |
| flight 2031 wors... | 2.0        | 2     |
| has the worst cu... | 4.0        | 4     |
| have time and th... | 4.0        | 4     |
| if it is ever co    | 4.0        | 4     |

01

1. Tune the hyperparameters of model.

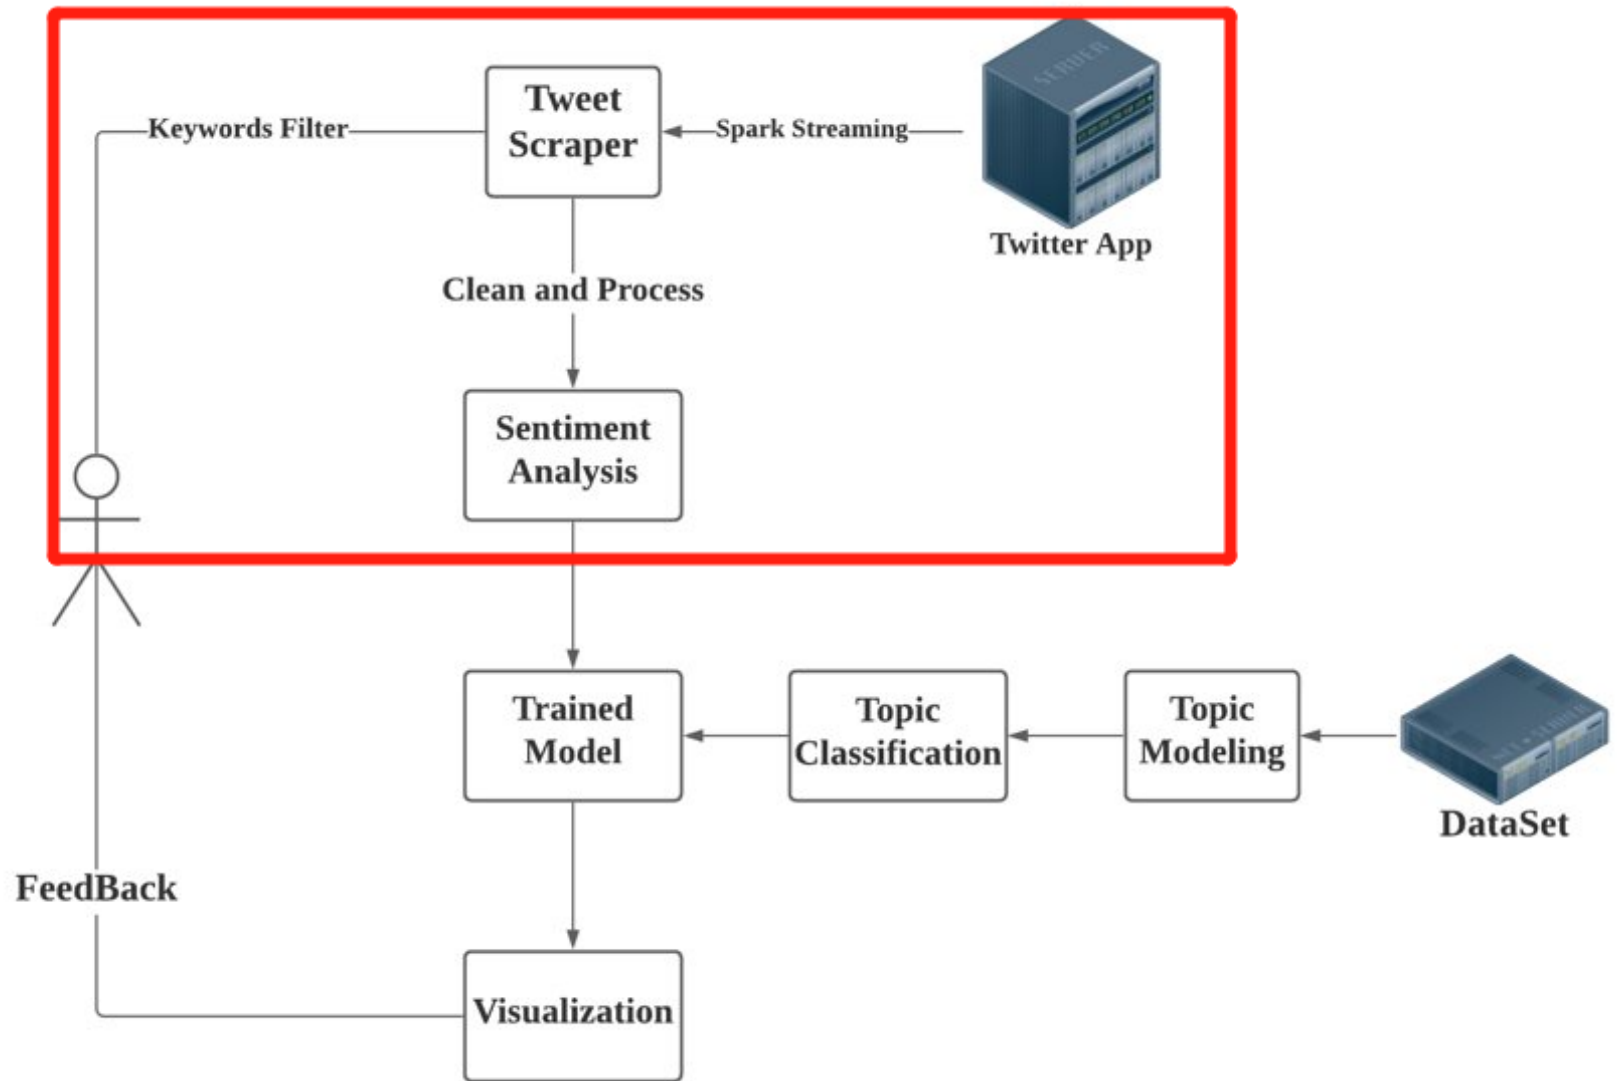
02

2. Make data more accurate.


03

3. Training with more data...

# Twitter Real-time Data



# Twitter Scraper



# Twitter Scraper

```
object TweetScraper {  
  def setupLogging(): Unit = {  
    import org.apache.log4j.{Level, Logger}  
    val rootLogger = Logger.getRootLogger  
    rootLogger.setLevel(Level.ERROR)  
  }  
  
  /** Configures Twitter service credentials using twitter.txt  
  def setupTwitter(): Unit = {  
    import scala.io.Source  
  
    val lines = Source.fromFile("data/actualdata/twitter.txt")  
    for (line <- lines.getLines) {  
      val fields = line.split(regex = " ")  
      if (fields.length == 2) {  
        System.setProperty("twitter4j.oauth." + fields(0), fields(1))  
      }  
    }  
    lines.close()  
  }  
}
```

```
def main(args: Array[String]) {  
  
  // Configure Twitter credentials using twitter.txt  
  setupTwitter()  
  
  // Set up a Spark streaming context named "PopularHashtags" that runs locally using  
  // all CPU cores and one-second batches of data  
  val ssc = new StreamingContext(master = "local[*]", appName = "PopularHashtags", Seconds(1))  
  
  // Get rid of log spam (should be called after the context is set up)  
  setupLogging()  
  
  // Create a DStream from Twitter using our streaming context  
  val keywords = Configure.tweetfiltersc.getString(path = "KEYWORDS").split(regex = ",").toSeq  
  println(keywords)  
  val tweets = TwitterUtils.createStream(ssc, None, keywords)  
  // Now extract the text of each status update into DStreams using map()  
  val statuses: DStream[String] = tweets.filter(t => t.getLang() == "en").map(status => status.getText)  
  
  val spark = SparkSession.builder  
    .master(Configure.sparkc.getString(path = "MASTER_URL"))  
    .appName(name = "TweetStream")  
    .getOrCreate()  
  
  spark.sparkContext.setLogLevel("ERROR")  
}
```

# Twitter Scraper with Kafka

---

```
def main(args: Array[String]): Unit = {
    // set log level
    setupLogging()


    val cb = new ConfigurationBuilder
    cb.setDebugEnabled(true)
        .setOAuthConsumerKey(Configure.twitter.getString( path = "CONSUMER_KEY"))
        .setOAuthConsumerSecret(Configure.twitter.getString( path = "CONSUMER_KEY_SECRET"))
        .setOAuthAccessToken(Configure.twitter.getString( path = "ACCESS_TOKEN"))
        .setOAuthAccessTokenSecret(Configure.twitter.getString( path = "ACCESS_TOKEN_SECRET"))
        .setJSONStoreEnabled(true)

    //create kafka props
    val props = new Properties()
    props.put("bootstrap.servers", Configure.kafkac.getString( path = "BOOTSTRAP_SERVERS"));
    props.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer"); // Str
    props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer"); // S

    val producer = new KafkaProducer[String, String](props)
    val kafkatopic = Configure.kafkac.getString( path = "TOPIC")
    val statuslistener = new StatusListener {
        /*
        StatusListener defines what to do with the tweets as they stream
        */
    }
```

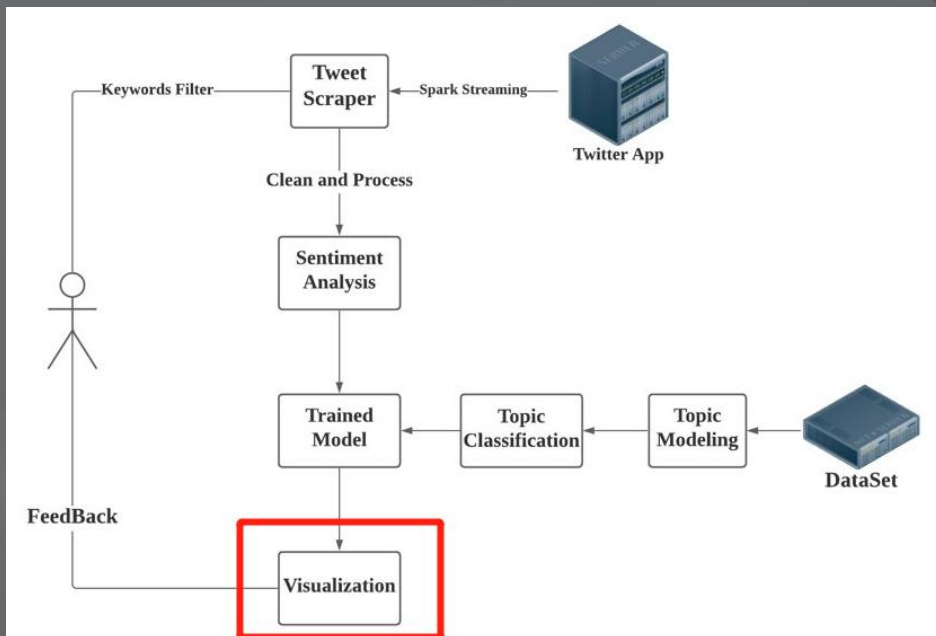
# Sentimental analysis with CoreNLP

```
object CleanTweets {  
  
  // Data from twitter would be like "RT @User: xxx", clean it  
  def clean(input: String): String = input.substring(input.indexOf(": ") + 2)  
  
}
```





# Upload to Elastic



```
/**  
 * Upload the data processed to Elastic Search  
 * @param dataframe Data whose type is dataframe  
 * @param path Elastic Search path  
 */  
def dataframeToElastic(dataFrame: DataFrame, path: String): Unit = {  
    dataFrame.write  
        .format( source = "org.elasticsearch.spark.sql")  
        .option("es.port", 9200)  
        .option("es.nodes", "localhost")  
        .mode( saveMode = "append")  
        .save(path)  
}
```

# Upload to Elastic

```
Text,senti,classify
"This is a test0",positive,0.0
"This is a test2",positive,2.0
```

GET localhost:9200/test0419/doc/\_search

Params Authorization Headers (7) Body Pre-request Script Tests Se

Body Cookies Headers (4) Test Results

Pretty Raw Preview Visualize JSON ↻

```
16 "hits": [
17   {
18     "_index": "test0419",
19     "_type": "doc",
20     "_id": "GWWA6HgBATphAarWfe0Y",
21     "_score": 1.0,
22     "_source": {
23       "Text": "This is a test0",
24       "senti": "positive",
25       "classify": "0.0"
26     }
27   },
28   {
29     "_index": "test0419",
30     "_type": "doc",
31     "_id": "GmWA6HgBATphAarWfe0Y",
32     "_score": 1.0,
33     "_source": {
34       "Text": "This is a test2",
35       "senti": "positive",
36       "classify": "2.0"
37     }
38   },
39 ]
```

Search

+ Add filter

tweetsairline

Search field names

Field filters 0

Records

Available fields 3

airline\_sentiment.keyword

prediction

text.keyword

Empty fields 0

Meta fields 3

Stacked bar

#

Drop some fields here to start

Kibana

Lens is a new tool for creating visualization

Make requests and give feedback

KQL

Last 4 days

Show dates

Refresh

Quick select

Last 4 days

Apply

Commonly used

Today

This week

Last 15 minutes

Last 30 minutes

Last 1 hour

Last 24 hours

Last 7 days

Last 30 days

Last 90 days

Last 1 year

Recently used date ranges

Last 4 days

This week

Last 2 days

Today

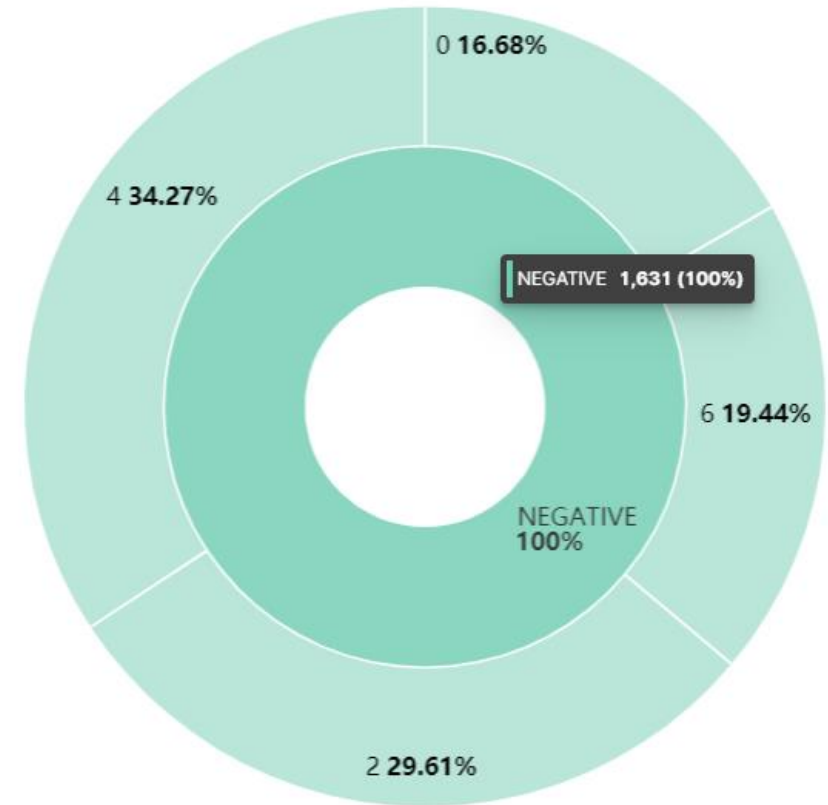
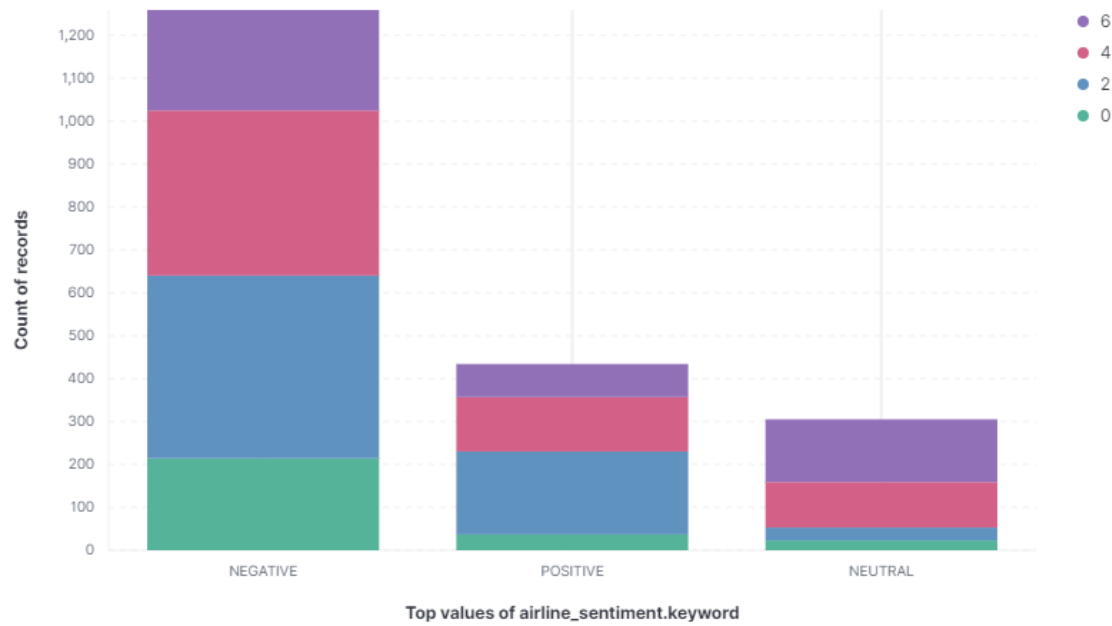
Last 24 hours

Refresh every

2 seconds

Start

# Example







Thank you for watching !